Connecting Databricks Genie Spaces to Glean

End-to-end setup guide and checklist for a multi-cloud (Azure + AWS) environment

---

1. Overview: the two layers

Your production data lives in Databricks workspaces across Azure and AWS. A Genie space is the "trained" semantic layer you build on top of that data — a curated set of tables plus business instructions and example queries that let people ask questions in natural language. To surface those spaces inside the Glean interface, you connect them in two complementary ways:

- **Native Databricks connector —** indexes your Genie spaces and AI/BI dashboards so they appear in Glean search (title, description, metadata), permission-trimmed. This is discovery: people can find the right space.
- **Databricks Tools / Genie action —** lets the Glean Assistant live-query a Genie space in natural language and return real answers from your data. This is the conversational "ask Glean a data question" experience.

Most people who ask to "connect trained Genie spaces to Glean" want the second layer; the first complements it. This guide sets up both. They are independent — you can do them in either order — but you want both for the full experience.

Note on multi-cloud: there is no single global configuration that spans Azure and AWS. You run the connector once per Databricks account and create one Genie tool per space, each pinned to its own host URL. Azure has one quirk throughout — it has no "list all workspaces" API, so you must always supply explicit workspace URLs.

2. Prerequisites

- ☐ Databricks Premium plan (required for API access used by the connector).
- ☐ AI/BI Dashboards enabled on the Databricks account.
- ☐ A running SQL Warehouse (serverless recommended) — Genie executes queries against a warehouse.
- ☐ Account Admin access to each Databricks account (Azure and AWS), or the ability to have a service principal made Workspace Admin on each workspace.
- ☐ Glean Admin Console access (to add connectors) and Glean platform/tools access (to add tools).

3. Phase 1 — Build and train the Genie spaces in Databricks

Do this in each cloud (Azure and AWS) where the underlying data lives. Genie spaces are cloud- and workspace-specific.

1. Get your tables into Unity Catalog and confirm a SQL Warehouse is running.
2. In the workspace, go to Genie → New, name the space, and add the tables and views it should reason over.
3. Train it: add general instructions (business context and definitions), example SQL queries for common questions, join hints, and column descriptions. Test in the Genie chat until answers are reliable.
4. Set permissions on the space for the users and groups who should reach it. Glean enforces these at query time — anyone who cannot see the space in Databricks will not get answers through Glean, so get this right here.
5. Record the identifiers below for each space — you will need them in Phases 2 and 3.

Value	Where to find it
Host URL	Browser address bar of the workspace. AWS: https://<instance>.cloud.databricks.com/ | Azure: https://adb-<instance>.<n>.azuredatabricks.net/
Genie Space ID	In the Genie space URL: .../genie/rooms/<genie-space-id>/...
SQL Warehouse ID	SQL Warehouses page in Databricks; the warehouse the space runs on.
Account URL	Full URL of the Databricks Account console after sign-in (per cloud).

4. Phase 2 — Index Genie spaces into Glean search (native connector)

This makes your Genie spaces and AI/BI dashboards discoverable in Glean search. Repeat the connector configuration once per Databricks account (one for AWS, one for Azure).

Step A — Prepare Databricks credentials

6. Sign in to the Databricks Account console for the cloud and copy the full Account URL from the address bar:

```
AWS:   https://accounts.cloud.databricks.com/login
Azure: https://accounts.azuredatabricks.net/login
```

7. Create a service principal named “Glean” (User management → Service principals → Add service principal).
8. Assign it the **Account Admin** role (recommended — needed to resolve group permissions across workspaces). If policy forbids that, use workspace-level mode instead and make the service principal **Workspace Admin** on every workspace you will index.
9. On the service principal → Credentials & secrets → OAuth secrets → Generate secret (720-day lifetime recommended). Copy the Client ID and the secret immediately — the secret is shown only once. This is an M2M OAuth flow.

Step B — Configure the connector in Glean

10. In the Glean Admin Console, go to Connectors → Add Connector → Databricks.
11. Enter the Account URL, OAuth Client ID, and OAuth Secret.
12. **For Azure only:** add every workspace explicitly (name, URL, and ID), because Azure has no list-workspaces API.
13. Click Save. Glean runs an initial full crawl and validates permissions. Your Genie spaces and dashboards then appear in Glean search, trimmed to each user’s Databricks permissions.

Heads-up: the connector uses full crawls only (no incremental sync or webhooks), so newly added spaces appear on the next full crawl rather than immediately.

Phase 2 checklist

☐ Account URL copied (per cloud).

☐ “Glean” service principal created with the correct admin role.

☐ OAuth Client ID + secret generated and stored.

☐ Connector added in Glean with the three credentials.

- ☐ Azure: all workspace URLs entered explicitly.
- ☐ Saved; initial crawl completed and spaces visible in search.

5. Phase 3 — Enable live Genie querying from the Glean Assistant (Tools)

This is the piece that lets Glean actually answer data questions by calling Genie. Create one tool per Genie space, since each carries its own space ID, warehouse ID, and host URL.

Step A — Create a Databricks OAuth app

Works the same for AWS and Azure. In the Databricks Account console → Settings → App Connections → Add Connection:

- Name it, e.g. “Glean Databricks Genie OAuth”.
- Add a dummy URL in the Redirect URL field for now.
- Set Access Scopes to All APIs and check Generate a Client Secret.
- Set token lifetimes (guidance: 6-hour access token, 3-month refresh token, so agents don’t re-authenticate often).
- Save, then copy the Client ID and Client Secret.

Step B — Configure the tool in Glean

14. In Glean, go to Settings → Platforms → Tools → Add → Databricks Tools.
15. Fill the Configuration section:

- Default Genie Space ID (from Phase 1).
- Default Warehouse ID.
- Host URL in the cloud-correct format:

```
AWS:   https://<instance>.cloud.databricks.com/
Azure: https://adb-<instance>.<n>.azuredatabricks.net/
```

16. Fill the Authenticate section:

- Client ID and Client Secret from the OAuth app.
- Authorization URL: https://<host-url>/oidc/v1/authorize
- Token URL: https://<host-url>/oidc/v1/token
- Scope: all-apis,offline_access

17. Save the tool. Glean generates a Callback URL — copy it back into the Databricks OAuth app’s Redirect URL, replacing the dummy one.
18. Add the tool to the Glean Assistant and/or the specific agents or prompts that should use it. Create a separate tool per Genie space and expose whichever ones each agent needs.
19. Test in the Assistant or agent builder: ask a natural-language question that should hit the space and confirm it returns real query results.

Phase 3 checklist

- ☐ Databricks OAuth app created (All APIs scope, client secret generated).
- ☐ Databricks Tools added in Glean with Genie Space ID, Warehouse ID, and Host URL.
- ☐ Authorization URL, Token URL, and scope set correctly.
- ☐ Tool saved; Callback URL copied back into the Databricks OAuth app.
- ☐ Tool added to the Assistant / relevant agents.

- ☐ Live query tested and returning results.
- ☐ Repeated per Genie space and per cloud as needed.

6. Things worth flagging

- **Multi-cloud means duplication.** Run the connector once per Databricks account and create one Genie tool per space, each pinned to its own host URL. There is no single config spanning Azure and AWS.
- **Permissions come from Databricks.** Access is enforced at query time, so anyone who cannot see a space in Databricks will not get answers through Glean. Manage access on the space itself.
- **The two phases are independent.** Search indexing (Phase 2) and live querying (Phase 3) can be set up in either order, but both together give the full experience.
- **Azure specifics.** Always provide explicit workspace URLs; the Azure host format is https://adb-<instance>.<n>.azuredatabricks.net/.

7. References

- Glean Databricks connector (indexing Genie spaces): docs.glean.com/connectors/native/databricks/
- Glean Databricks Tools setup (live Genie querying): docs.glean.com/administration/tools/setup-tools/databricks-tools-setup
- Connect Databricks Genie to Glean Assistant: docs.glean.com/administration/assistant/warehouse-data/connect-databricks-genie-to-glean-assistant
- Glean + Databricks Genie announcement: glean.com/blog/glean-databricks-genie-announce